



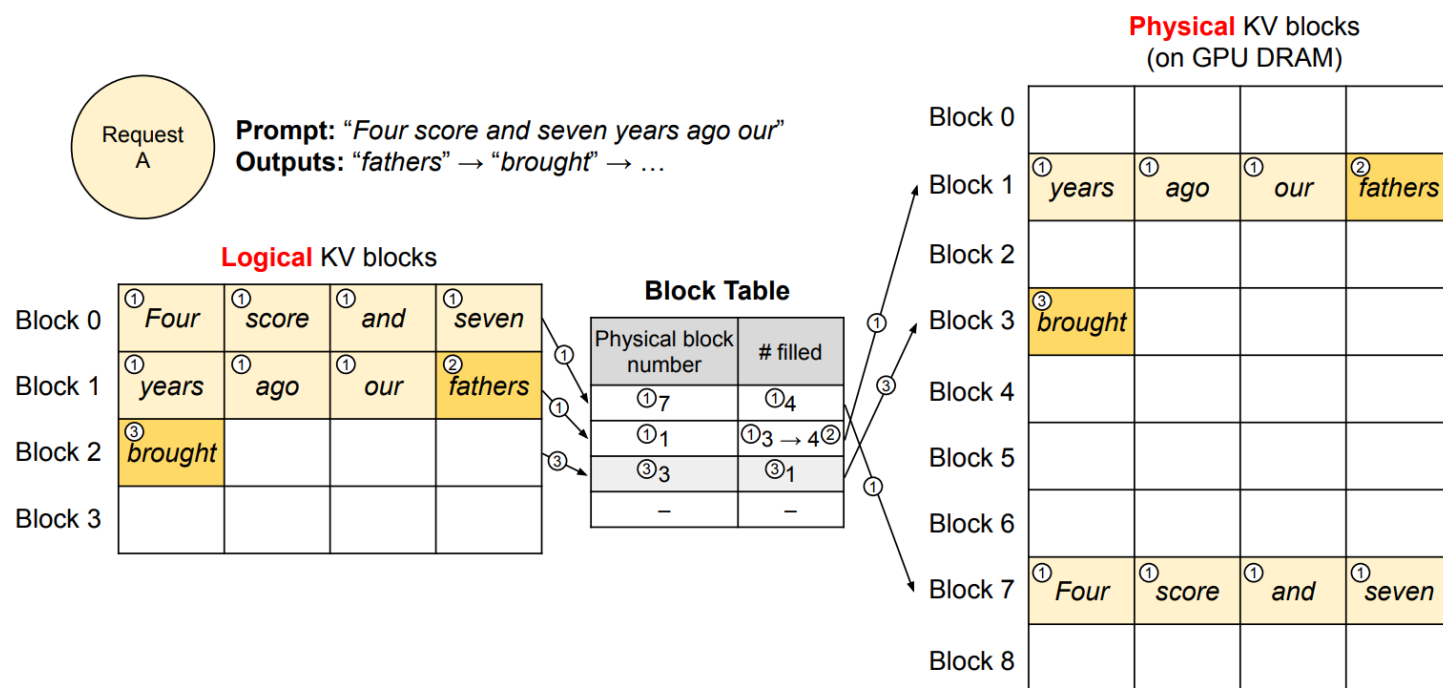
FlashInfer

INTRODUCTION TO LLM
INFERENCE SERVING SYSTEMS

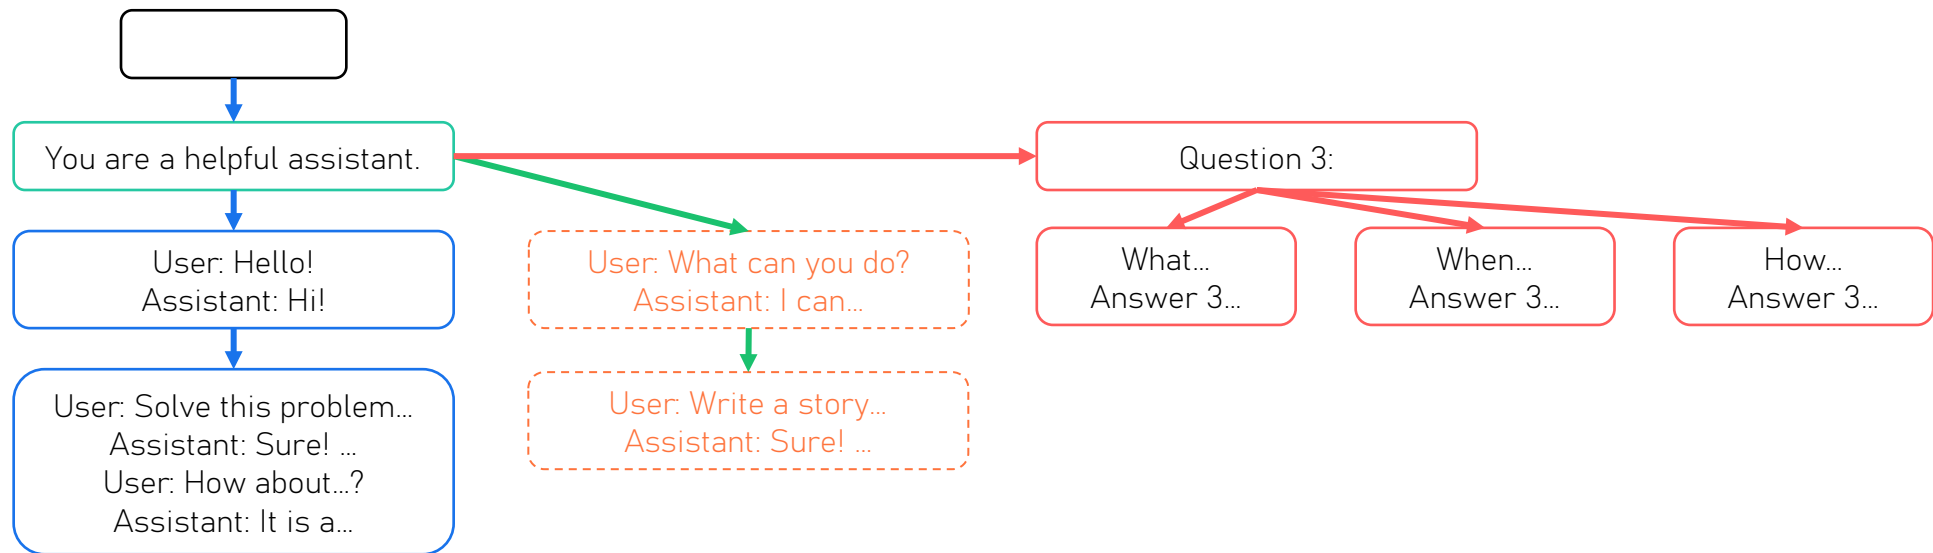
CHUHONG YUAN



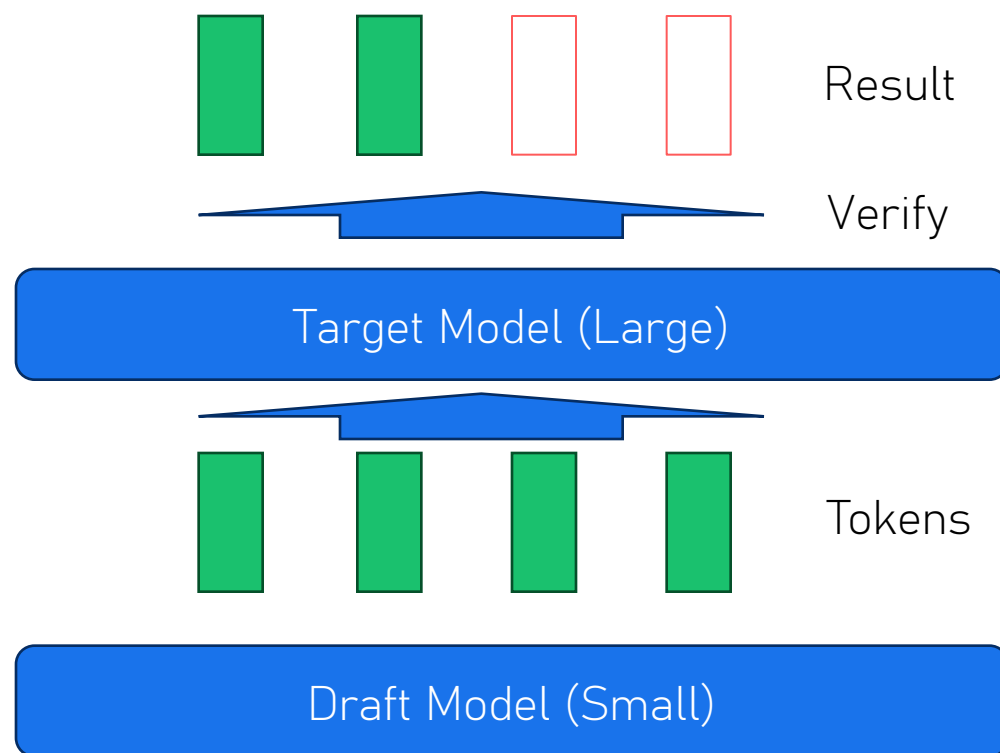
Review of PageAttention




Review of RadixAttention



Background – Speculative Decoding



Background – Speculative Decoding




Background – Sparse Attention

	K1	K2	K3
Q4	√	√	√
Q5	√	√	√
Q6	√	√	√

Background – Sparse Attention

	K1	K2	K3
Q4	√	√	√
Q5	√	√	√
Q6	√	√	√



	K1	K2	K3
Q4	√	√	
Q5		√	√
Q6	√		

Background – KV Cache Sparsity

	I	am	like	the
I	√			
am	√	√		
like	√		√	
the	√	√		√



Block-Sparse Matrix

- Still use a paged KV Cache management
- Allow sparsely selecting the K/V values to compute
- Organize the data into blocks to fit the size of tensor cores
- The block size can be customized




Block-Sparse Matrix - Usage






Block-Sparse Matrix - Usage

Row size = 4






Block-Sparse Matrix - Usage

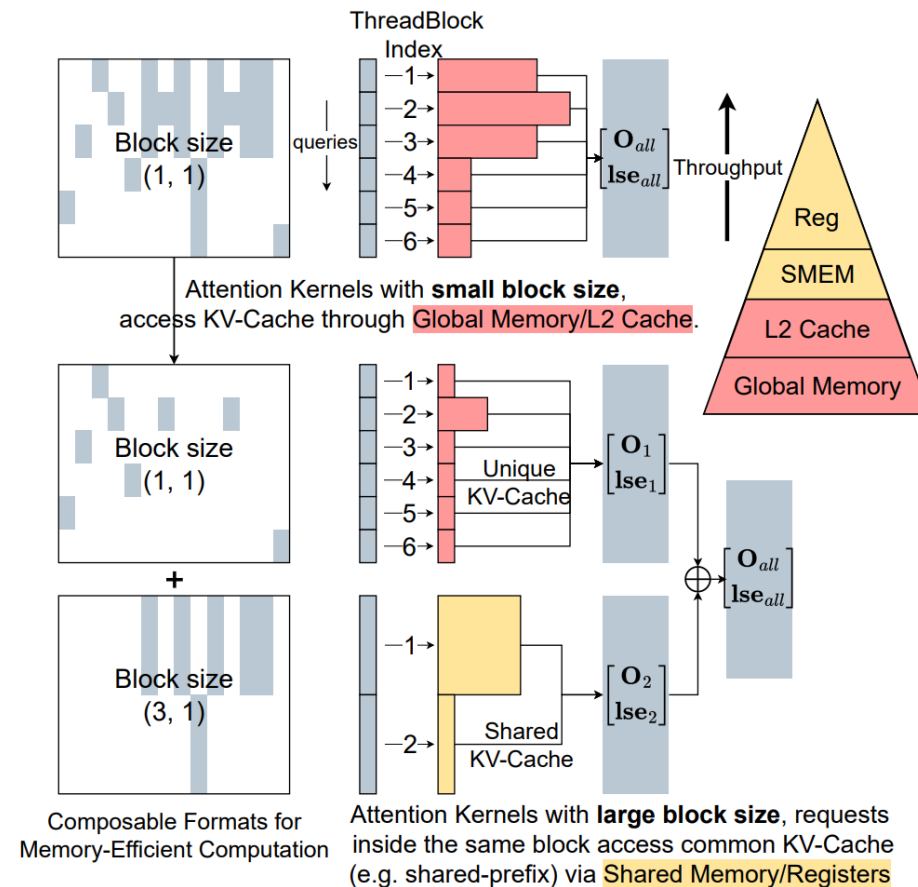




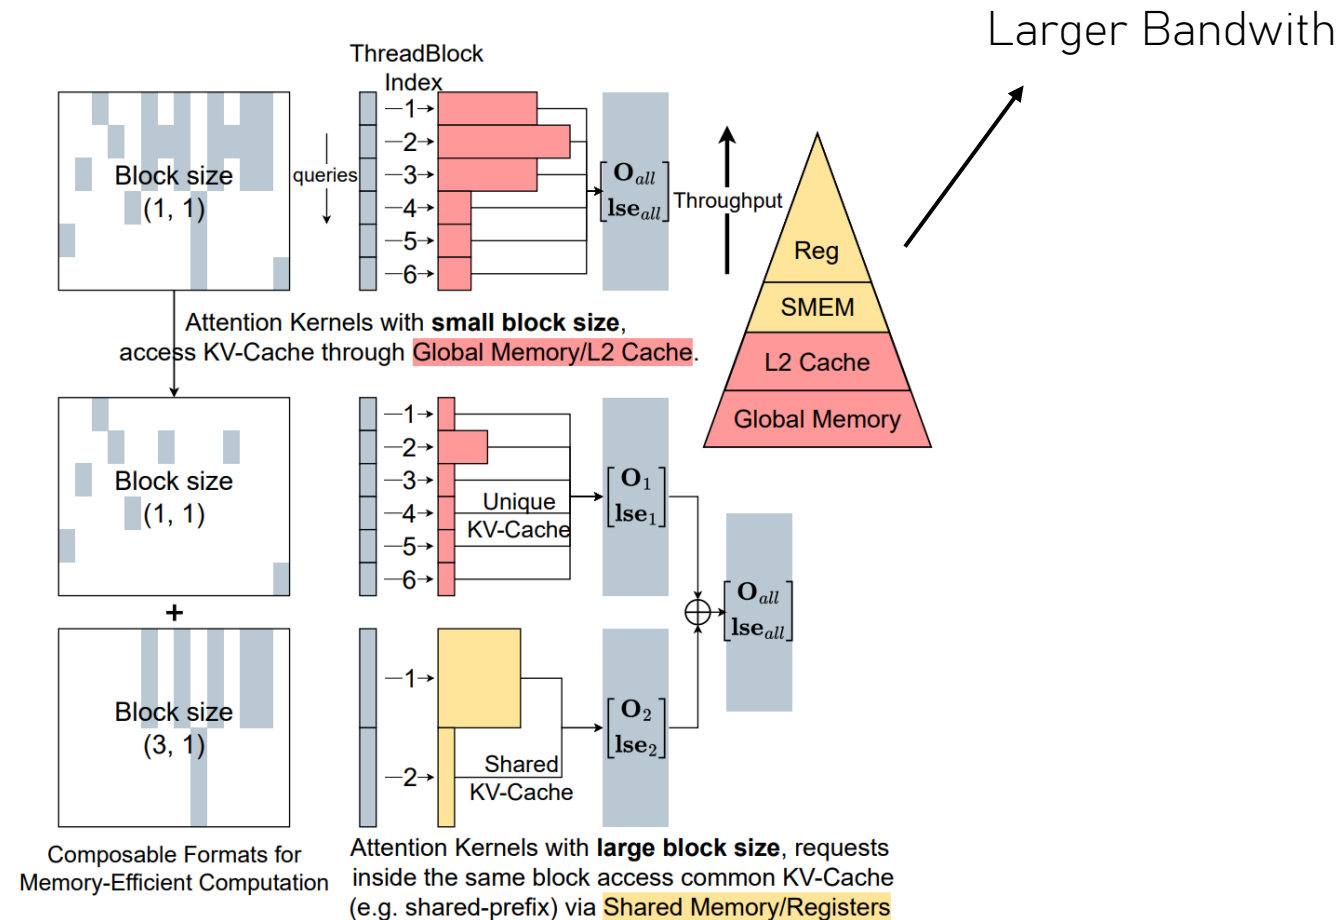
Block-Sparse Matrix - Usage



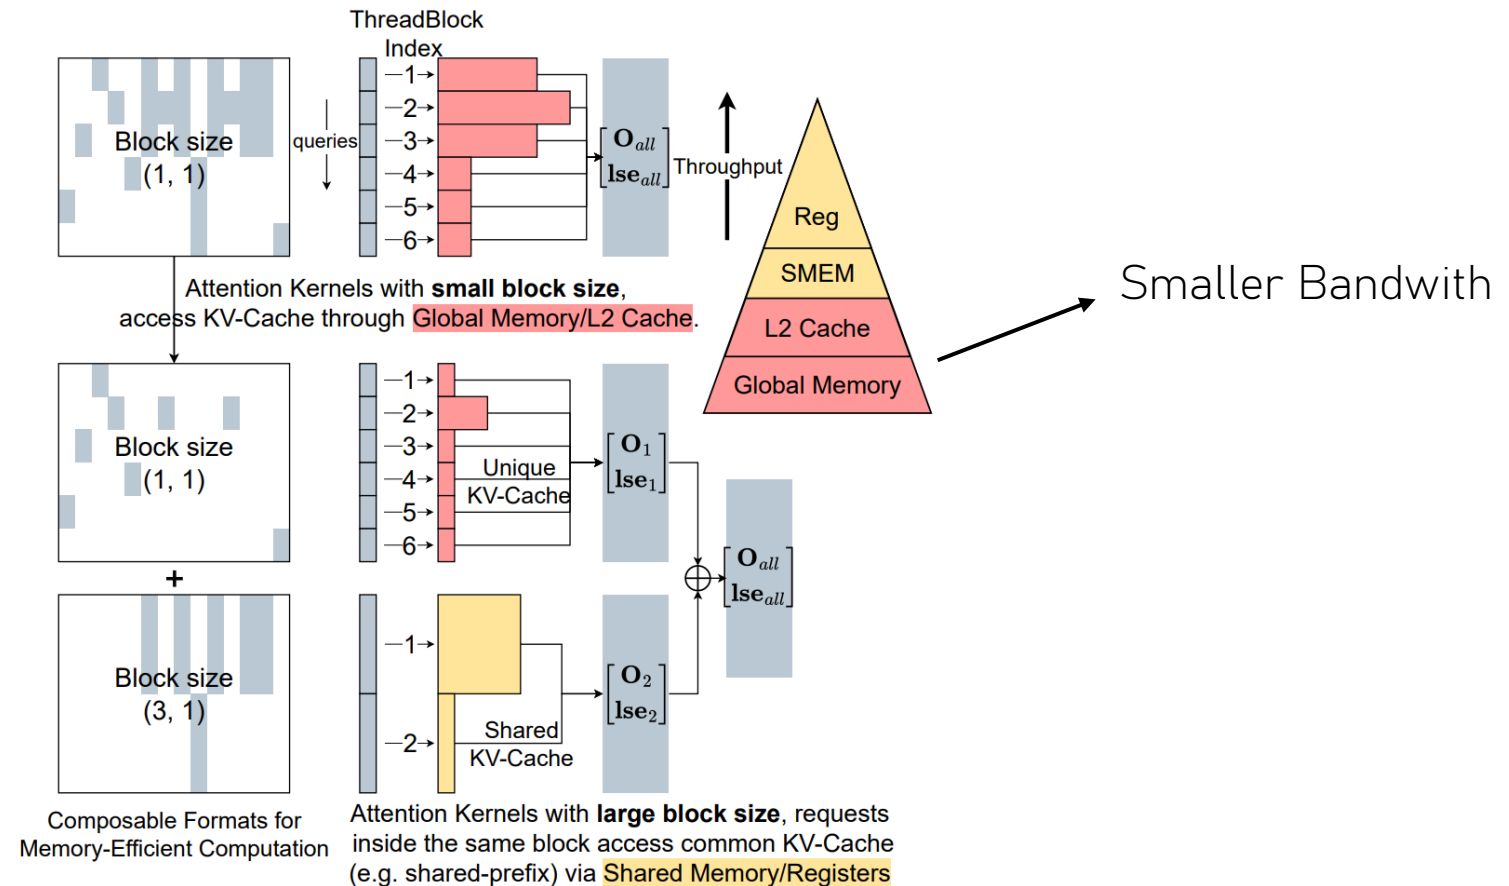
Composable Formats



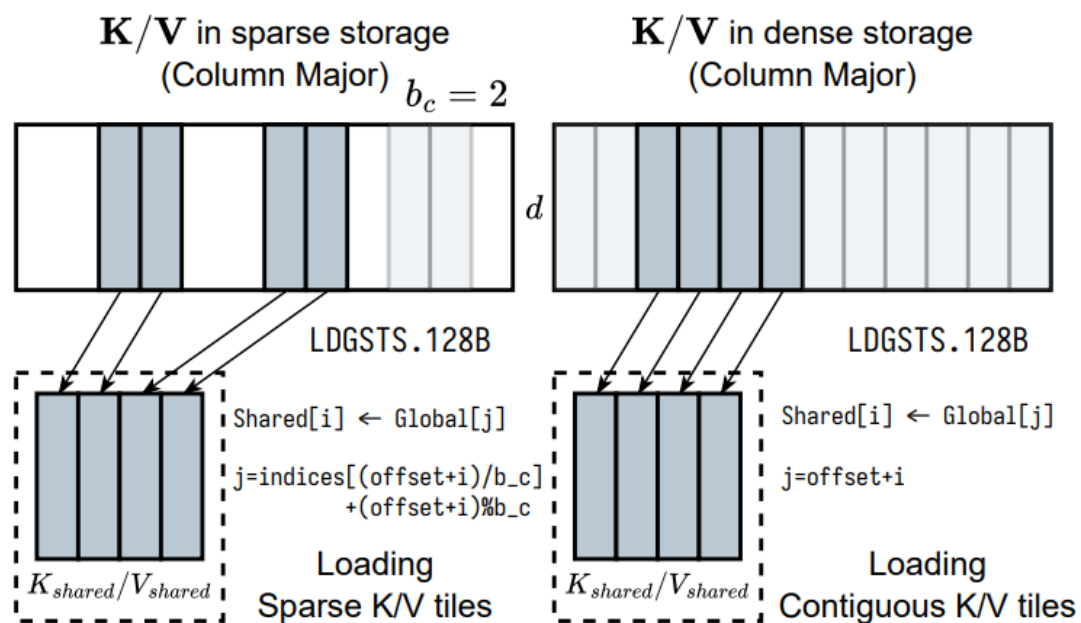
Composable Formats



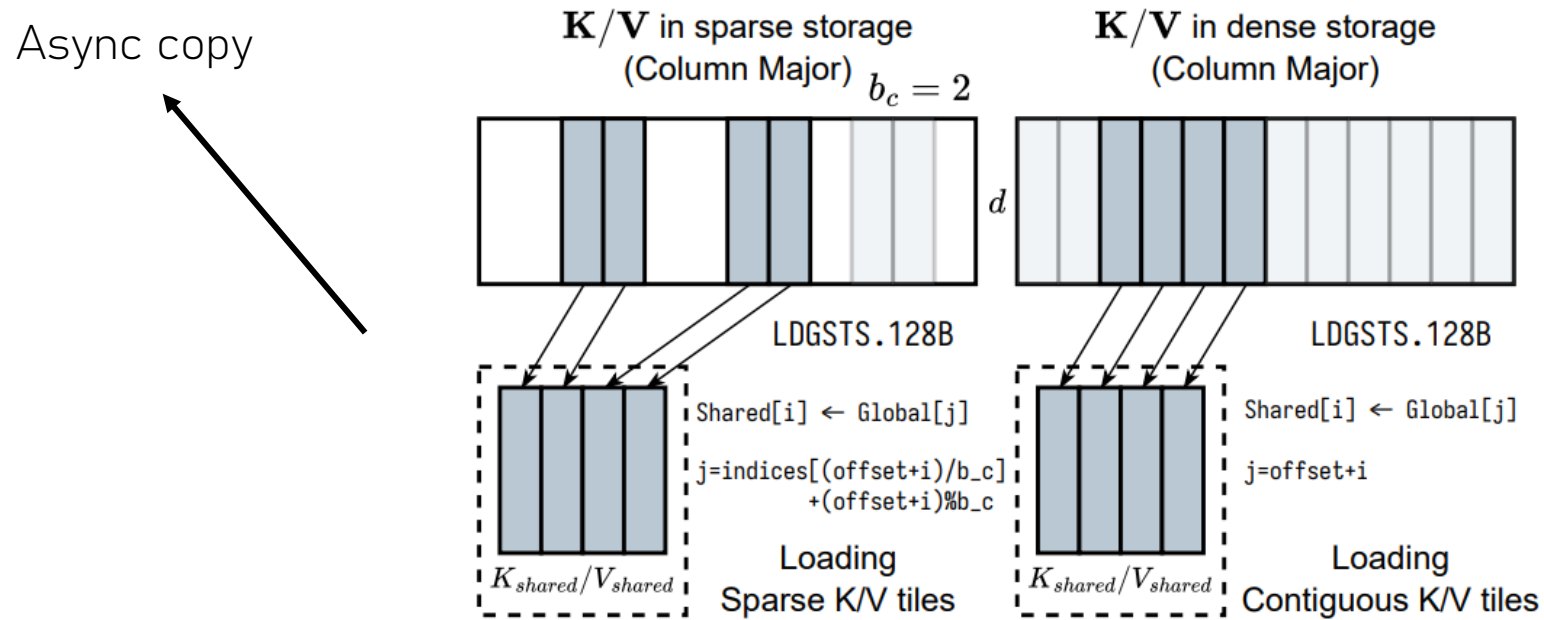
Composable Formats



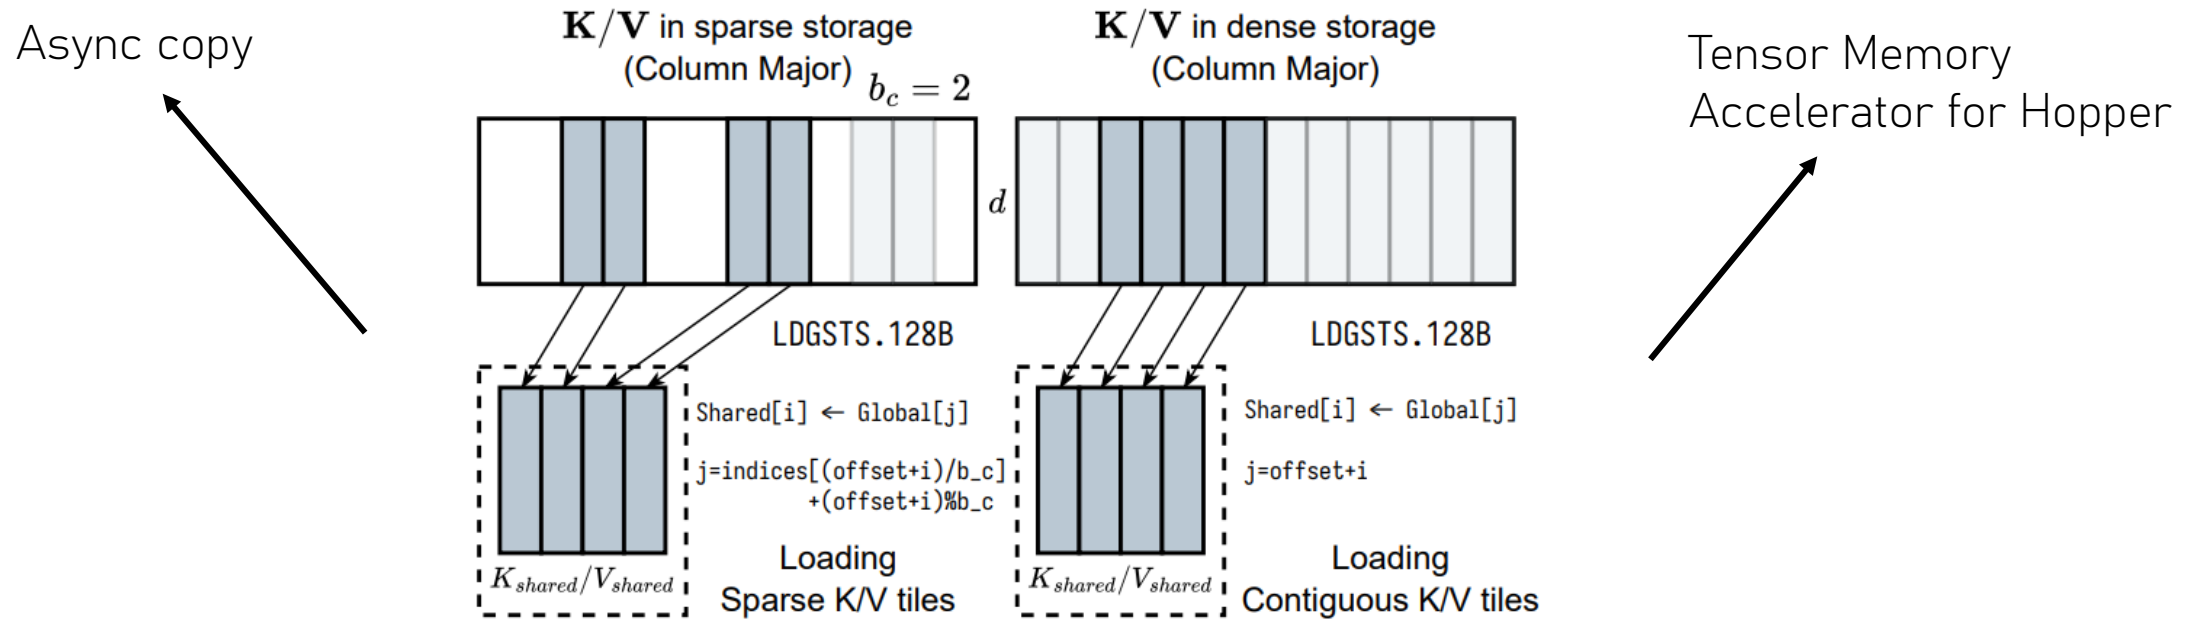
Global To Shared Memory Data Movement



Global To Shared Memory Data Movement




Global To Shared Memory Data Movement



JIT Compiler – Background

- Different attention algorithms
- Grouped-query attention
 - A group of query heads share the same K and V





JIT Compiler

- Enable defining the specifications of attention variants
- Provide interfaces that can be customized
 - QueryTransform, KeyTransform, ValueTransform
 - OutputTransform
 - LogitsTransform, LogitsMask

JIT Compiler – Example

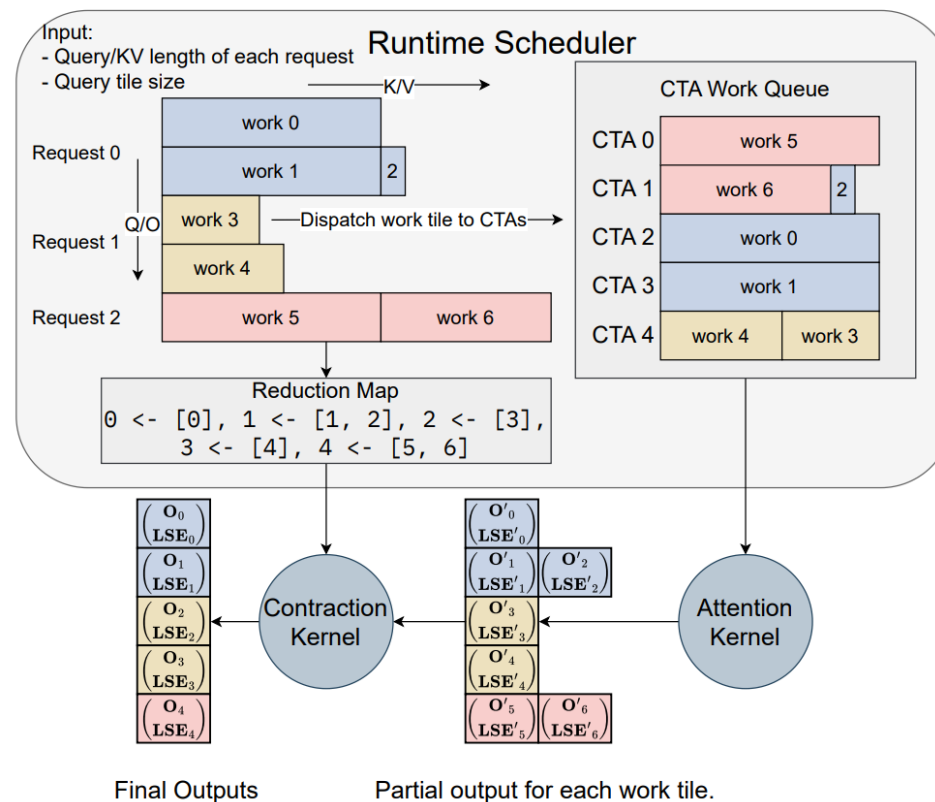
Attention Specification in Python	Part 1: Kernel Parameters Class	Part 2: Kernel Traits class
<pre>spec_decl = r""" template <typename Params_, typename KernelTraits_> struct FlashSigmoid { using Params = typename Params_; using KernelTraits = typename KernelTraits_; static constexpr bool use_softmax = false; float scale, bias; FlashSigmoid(const Params& params, int batch_idx, uint8_t* smem_ptr) { // Copy from CUDA constant memory to registers scale = params.scale; bias = params.bias; } ... float LogitsTransform(const Params& params, float logit_score, int batch_idx, int qo_idx, int kv_idx, int qo_head_idx, int kv_head_idx) { return 1. / (1. + expf(-(logits_score * scale + bias))); } }; """ attn_spec = AttentionSpec("FlashSigmoid", dtype_q, dtype_kv, dtype_o, idtype, head_dim, is_sparse, additional_vars=[("scale", "float"), ("bias", "float")], additional_tensors=[], spec_decl=spec_decl)</pre>	<pre>template <typename DTypeQ, typename DTypeKV, typename DTypeO, typename IdType> struct Params { DTypeQ* q; DTypeKV* k, v; DTypeO* o; float* lse; IdType* qo_indptr, kv_indptr, kv_indices, kv_seq_lens; ... // (generated) additional vars float scale; float bias; }; Part 4: Register custom operators in PyTorch torch::Tensor attention_call(torch::Tensor q, torch::Tensor k, torch::Tensor v, ... float scale, float bias // (generated) additional vars) { ... auto kernel = KernelTemplate<FlashSigmoid<Params<{{dtype_q}}, {{dtype_kv}}, {{dtype_o}}, {{idtype}}>>, KernelTraits>>; ... } // Register torch custom ops TORCH_LIBRARY_IMPL("FlashSigmoid", CUDA, m) { m.impl("run", &attention_call); }</pre>	<pre>struct KernelTraits { static constexpr HEAD_DIM = {{head_dim}}; static constexpr IS_SPARSE = {{is_sparse}}; ... }; Part 3: Kernel Body template <typename AttentionSpec> __global__ KernelTemplate(AttentionSpec::Params params) { ... // Init attention specification class. AttentionSpec attn(params, batch_idx, smem_ptr); ... // Iterate over all elements inside the thread logits tile. for (int i = 0; i < size(logits_tile); ++i) { // convert register index i to qo_idx, kv_idx, etc. qo_idx = get<0>(logits_tile(i)); kv_idx = get<1>(logits_tile(i)); ... logits_tile(i) = attn.LogitsTransform(params, logits_tile(i), batch_idx, qo_idx, kv_idx, qo_head_idx, kv_head_idx); } ... }</pre>



Load-Balanced Scheduling

- Motivation: different categories of workloads
 - Prefill – initial input without the need to do attention with previous tokens
 - Decoding – only a query vector each time
 - Append – additional inputs that need to do attention with previous tokens
- Schedule based on the needed KV Cache size for one chunk

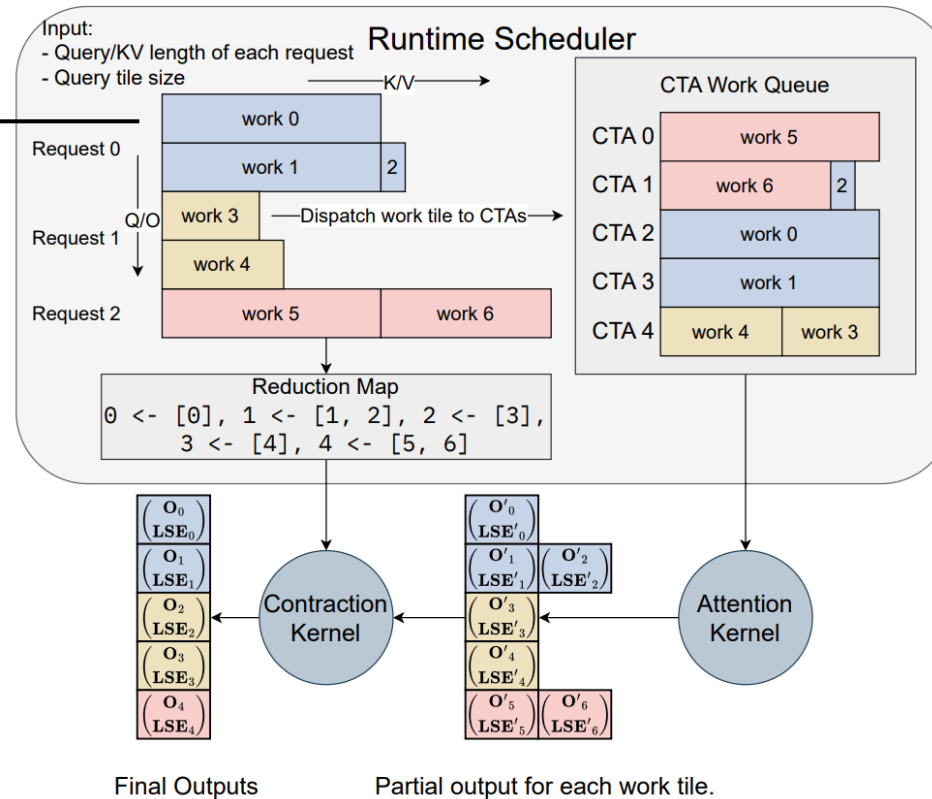
Load-Balanced Scheduling



Load-Balanced Scheduling

Length means cost

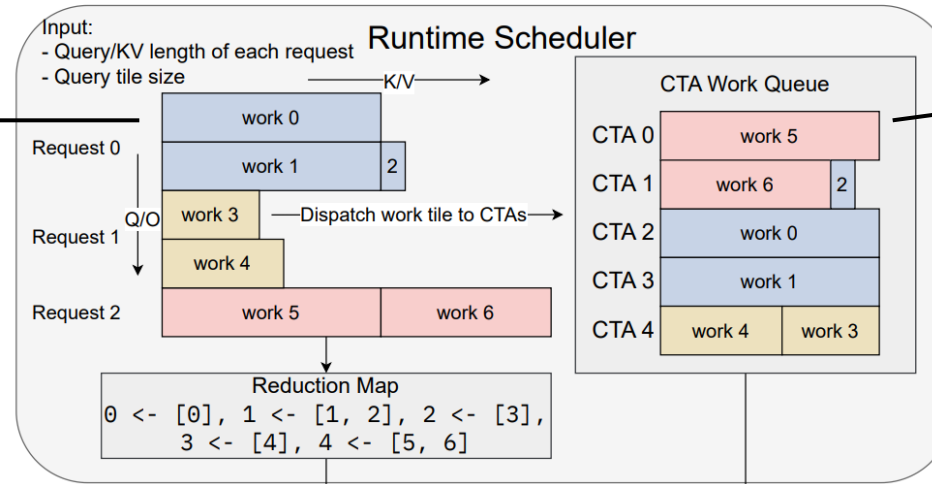
Cost is calculated by the length of Q, K, V




Load-Balanced Scheduling

Length means cost


Cost is calculated by the length of Q, K, V



Assign works with a priority queue



Final Outputs



Partial output for each work tile.

Contraction Kernel

Attention Kernel



Evaluation – Experiment Settings

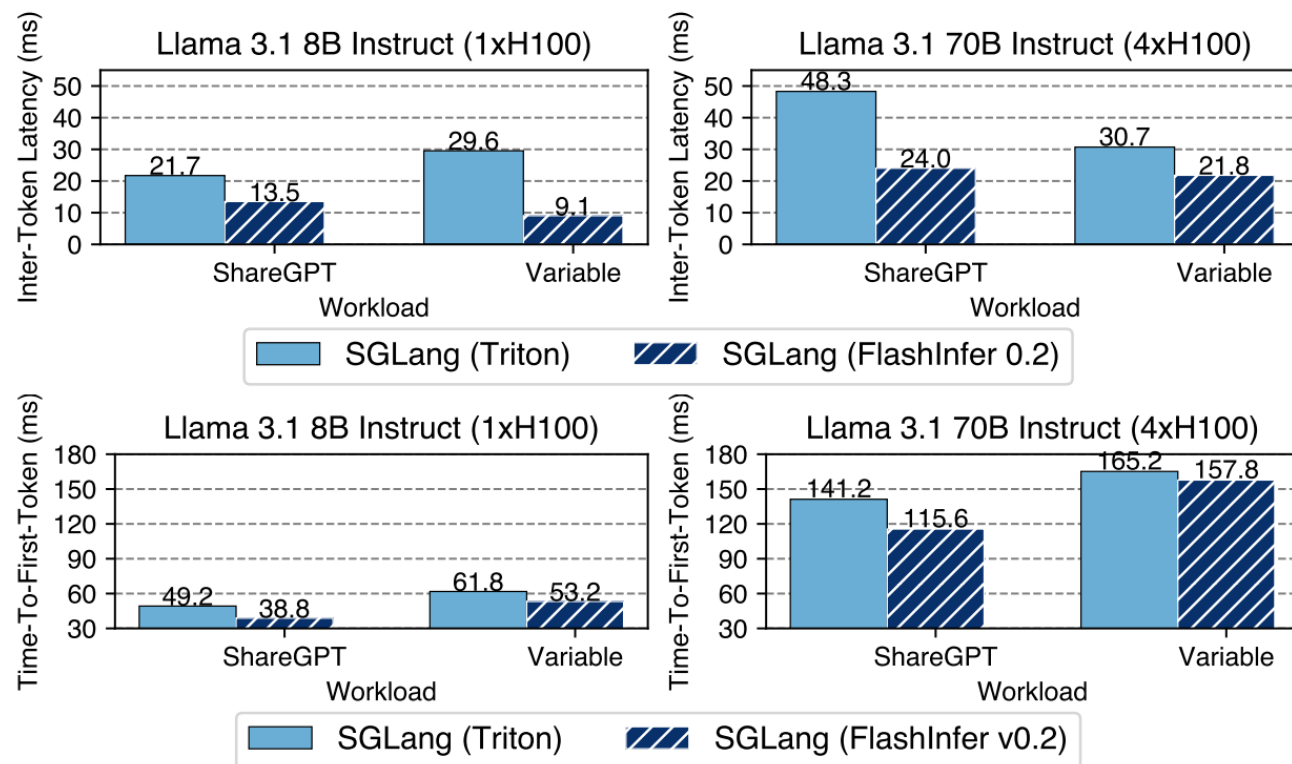
- NVIDIA A100 40GB, H100 80GB
- CUDA 12.4
- Pytorch 2.4.0
- Precision: f16
- Serving engine: SGLang 0.3.4
- Models: Llama3.1 8B, 70B, Vicuna-13B (fine-tuned from Llama)



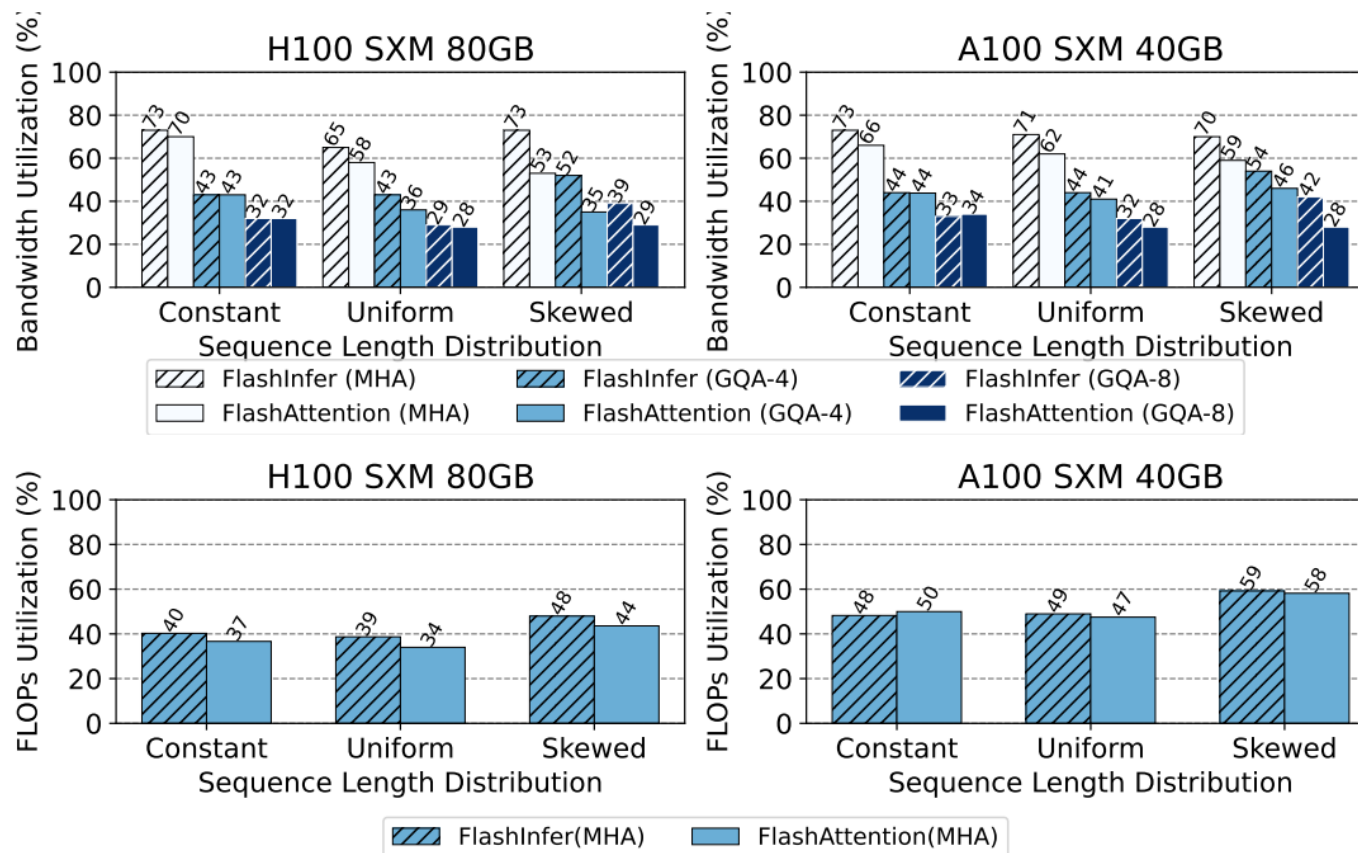
Evaluation – Workloads & Metrics

- Datasets
 - ShareGPT
 - Variable: different sequence lengths in different distribution (512-2048)
- Metrics
 - TTFT and TPOT (median)
 - Bandwidth and FLOPs utilization

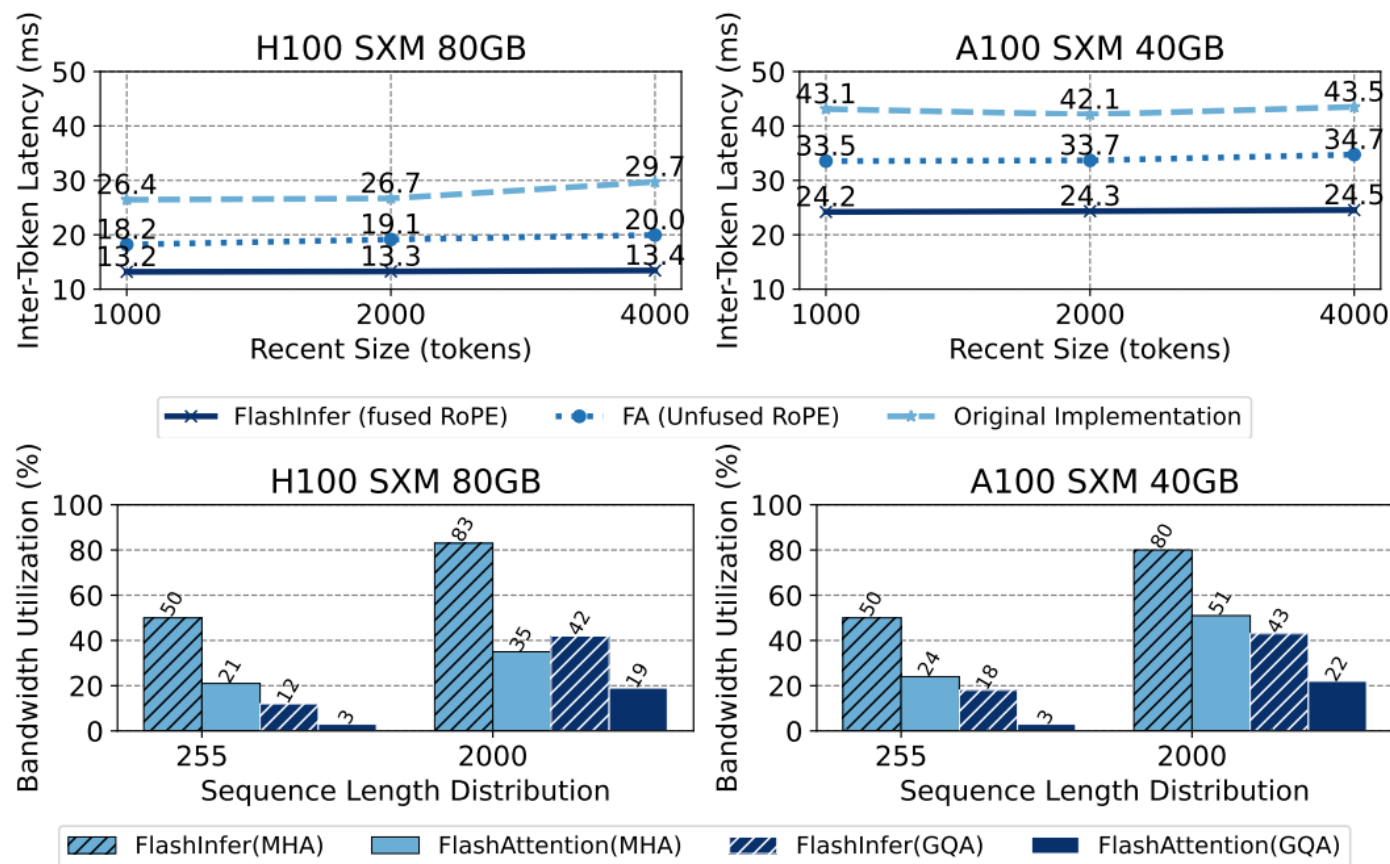
Evaluation – Latency



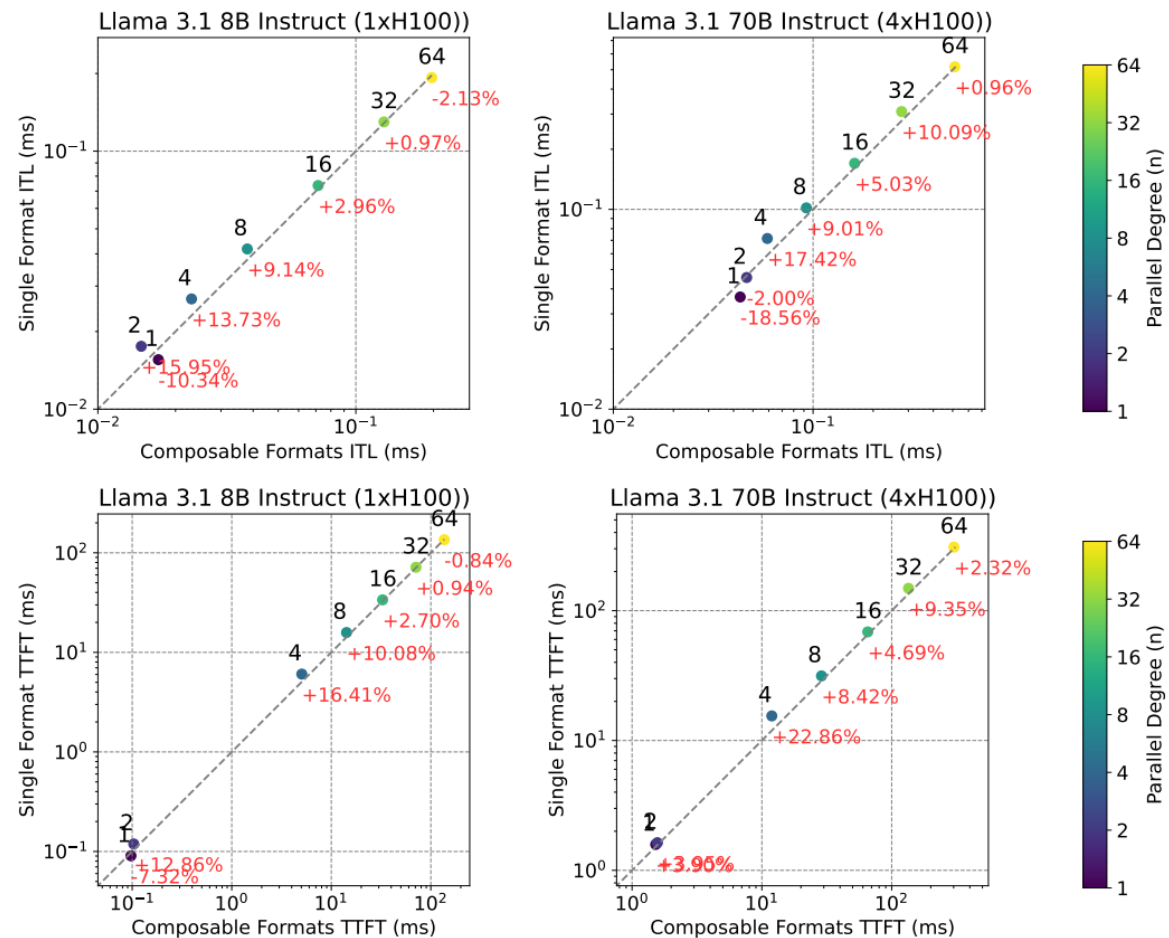
Evaluation – Kernel Performance



Evaluation – Long-Context Inference



Evaluation – Parallel Generation





Homework

- Read the paper **Taming Throughput-Latency Tradeoff in LLM Inference with Sarathi-Serve**, summarize the paper, specifically, including the points below:
 - What are the motivations/challenges of this work?
 - How does the design of this paper address the challenges?
 - How does the paper evaluate its design (experiment settings, workloads, metrics)?
 - How does the evaluation prove its claims?
- Related link:
 - Paper of Sarathi-Serve: <https://www.usenix.org/system/files/osdi24-agrawal.pdf>[Links to an external site.](#)

Q & A

